

gheim: A Swiss-Tuned PII Detector and Companion Dataset for the Four Official Languages of Switzerland and English

Joel Barmettler
Independent
Zürich, Switzerland
jbarmettler@proton.me

ABSTRACT

We release gheim-ch-pii-171k, a 171,336-chunk personally-identifiable information (PII) named-entity-recognition dataset covering the four official Swiss languages (Swiss German, Swiss French, Swiss Italian, Romansh) and English, and gheim-ch-560m, a 560M-parameter xlm-roberta-large fine-tune released alongside it. The dataset combines real Swiss text from court rulings, federal parliament records, Swiss-filtered web text and a Romansh corpus with an LLM-generated synthetic gap-fill layer that covers cells where the real-text corpus was too sparse for training and evaluation. The model attains a strict-span F_1 of 0.916 on a held-out test split, outperforming the next-strongest open contestant (openai/privacy-filter zero-shot) by 47 percentage points on the same Swiss test set, and reaching 0.89–0.96 person-NER F_1 on four widely-used external benchmarks. The dataset ships under CC BY 4.0; the model under Apache 2.0; both are available on the Hugging Face Hub. Reproducible eval scripts and checkpoints are on GitHub.

1 INTRODUCTION

The dominant pattern in production large-language-model (LLM) workflows is to call a hosted API: OpenAI, Anthropic, Google, Mistral. For Swiss-market deployments this raises a recurring compliance concern: every request that contains a customer name, a Swiss address, an IBAN, or an AHV/AVS social-security number leaves Swiss jurisdiction the moment it crosses the API boundary. The legal landscape for using third-party LLM services with personal data under Swiss law is analysed in detail by Rosenthal and Veraldi [19]; the practical mitigation in production systems is to anonymise sensitive content before the request and restore the originals on the response. The accuracy of the round-trip depends on the quality of the underlying named-entity recogniser.

Existing public PII detectors handle Swiss text poorly. The closest open generalist, openai/privacy-filter [15] (an OLMoE-based [13] 1.4B-parameter mixture-of-experts model with 50M active parameters), reaches a strict-span F_1 of 0.44 on our held-out Swiss test set despite multilingual training. Microsoft Presidio [12], configured with per-language NLP engines, reaches 0.43. Multilingual NER baselines such as Davlan’s xlm-roberta-base-ner-hrl [1] reach 0.65 person-only F_1 . None of the available systems fluently handle the four official Swiss languages together with the structured PII formats common in Swiss documents (CH-IBAN, AHV, VAT-CHE), and only one (SwissNER [24]) provides a Romansh evaluation slice.

This work contributes:

- (1) gheim-ch-pii-171k, a multilingual PII NER dataset of 171,336 chunks across Swiss German, Swiss French, Swiss

Italian, Romansh and English, covering eight PII categories (account number, address, date, email, person, phone, URL, secret), encoded in a 33-class BIOES tag scheme byte-aligned with openai/privacy-filter.

- (2) gheim-ch-560m, a 560M-parameter token-classification model fine-tuned from xlm-roberta-large [5] on the dataset, attaining a held-out test F_1 of 0.916 strict-span and 0.958 character-level.
- (3) A reproducible evaluation harness that scores both the released model on four established external benchmarks (ai4privacy/openpii-1m, ZurichNLP swissner, Babelscape WikiNeural [20], CoNLL-2003 [22]) and seven competing open detectors on the released test set, with replication checks against each baseline’s published numbers on its native dataset.

The rest of the paper is structured as follows. Section 2 situates the work against prior art. Section 3 describes the dataset construction. Section 4 covers model selection and training. Section 5 reports the evaluation, including the methodology-validation cross-checks. Section 6 discusses limitations and Section 7 concludes.

2 RELATED WORK

2.1 PII detection and de-identification

The dominant industrial open-source PII detector is Microsoft Presidio [12]: a per-language spaCy NER engine combined with checksum-validated regex recognisers for structured formats (IBAN, credit card, SSN, etc.). Presidio’s strength is its large built-in catalogue and easy extension; its weakness on Swiss text is reproducible from the per-spaCy NER backbone (Section 5, Table 6).

The AI4Privacy family of corpora and fine-tuned checkpoints [2, 3] is the most prolific recent open-source effort. Their datasets are templated synthetic prose with embedded Faker-generated PII at the million-chunk scale across 12+ fine-grained categories and 17+ locale variants, and the published distilbert fine-tunes (e.g. Isotonic/distil-bert_finetuned_ai4privacy_v2) are widely-downloaded open PII detectors on the Hugging Face Hub. Section 5 confirms the published per-token accuracy numbers but shows a sharp drop on real Swiss text: when scored under strict-span F_1 on gheim-ch-pii-171k the fine-tunes reach 0.16. The cross-domain gap reflects template-vs-real domain shift more than label-noise: AI4Privacy’s underlying templates were not designed with Swiss surface forms (CH-IBAN, AHV/AVS, +41 phones, German-language addresses) as targets.

openai/privacy-filter [15] is the strongest open zero-shot generalist on our test set. It is a token-classification head atop OLMoE [13], a 1.4B-total-parameter mixture-of-experts encoder with

50M parameters active per token, and it shares our 33-class BIOES schema. Its `openai-namespaces` release made it the default upstream from which our schema is byte-aligned, so a fine-tune started from this model can reuse the pretrained classifier-head weights.

2.2 Multilingual NER

XLM-RoBERTa [5] is the de-facto starting point for multilingual token-classification and is the base of `gheim-ch-560m`. The Davlan family of fine-tunes [1] are trained on PER/ORG/LOC across many European languages and provide our generalist baseline. WikiNeural [20] provides large-scale silver-data multilingual NER training and evaluation across nine European languages and is one of our cross-domain external benchmarks. The classic CoNLL-2003 shared task [22] remains the reference English NER benchmark and is also part of our cross-domain evaluation. The BIOES tag scheme used throughout follows the design analysis of Ratnov and Roth [18].

2.3 Swiss-NLP and the Apertus stack

SwissBERT [23] is the prior state-of-the-art encoder pretrained on Swiss text covering all four official languages, including Romansh. We considered SwissBERT as a base in the bake-off (Section 4, Table 3) and selected XLM-R-large by 0.8 pp validation F_1 ; SwissBERT remains a competitive alternative when ~ 300 M parameters is the deployment ceiling. Its companion test set, SwissNER [24], is the only public Swiss-news NER set that includes Romansh and is one of our cross-domain external benchmarks. The Romansh evaluation in particular relates closely to the WMT24++ benchmark extension of Vamvas et al. [25], which adds Rumantsch Grischun together with the five regional Romansh idioms (Sursilvan, Sutsilvan, Surmiran, Puter, Vallader) to the WMT machine-translation benchmark; the Apertus Romansh corpus, which underlies our `rm` chunks, is constructed in a similar idiom-aware way.

The real-text portion of our dataset is sourced from the Apertus pretrain corpora released alongside the Apertus Swiss-language LLM [6]. Apertus aggregates four Swiss-domain sources: the Open Legal Data Platform *Entscheidungsuche*, which provides machine-readable Swiss court decisions across cantons in their original German, French, or Italian [4]; *Curia Vista*, the official business database of the Swiss Federal Assembly, which includes parliamentary motions, postulates, and transcripts in German, French, and Italian [21]; the Swiss-filtered subset of FineWeb-2 [17] (the multilingual web-crawl corpus that scales the FineWeb pipeline of Penedo et al. [16] to all languages and applies XLM-R-based quality filters); and a purpose-built Romansh corpus drawing on cantonal law texts, the Lia Rumantscha online dictionaries, the ZurichNLP bilingual corpus, and Romansh Wikipedia. These corpora are released under CC-BY-4.0-compatible licences and are the substrate for the real-Swiss-text portion of `gheim-ch-pii-171k`.

2.4 Quantisation for browser deployment

The released model ships an `int8` ONNX export so the demo at `gheim.ch` runs entirely in-browser via WebGPU or WebAssembly. The quantisation procedure follows the post-training integer-arithmetic-only inference framework of Jacob et al. [9], applied

through ONNX Runtime [11] via the optimum dynamic-quantizer with AVX-512 VNNI as the deployment target. The browser runtime is `Transformers.js` [8].

2.5 Industry tooling

`transformers` [26] is the training and inference runtime. `seqeval` [14] provides the strict-span scoring used both for our headline F_1 and for every comparison and replication entry in Section 5. `spaCy` [7] provides the per-language NER pipelines used in our comparison and inside *Presidio*. `AdamW` [10] is the optimiser. `Faker_CH` provides the Swiss-locale random data used in our synthetic gap-fill payload generation.

3 DATASET: GHEIM-CH-PII-171K

3.1 Source corpora and language coverage

Real-text content is drawn from the Apertus pretrain corpora [6]: the Swiss legal-decisions subset *entscheidungsuche* curated from the Open Legal Data Platform of the same name [4], federal-parliament proceedings drawn from the *curia_vista* business database [21], the Swiss-filtered slice of FineWeb-2 [17], and the Romansh corpus released alongside Apertus, which compiles material from cantonal law texts, the Lia Rumantscha online dictionaries, the ZurichNLP bilingual corpus, and Romansh Wikipedia. Source documents are chunked at sentence boundaries to roughly 1,500 characters, then tagged with `facebook/fasttext-language-identification` at confidence ≥ 0.5 . Chunks below threshold are dropped. The five-language distribution of the released dataset is shown in Table 1.

Table 1: Per-language chunk distribution.

Code	Variant	Chunks	%
<code>de_ch</code>	Swiss German (written std.)	48,564	28.3
<code>fr_ch</code>	Swiss French	48,097	28.1
<code>it_ch</code>	Swiss Italian	45,462	26.5
<code>rm</code>	Romansh	17,171	10.0
<code>en</code>	English (Swiss-context)	12,042	7.0
Total		171,336	100

Spoken-form Swiss German (GSW) is not represented: the corpus is written-standard German as used in Swiss legal and journalistic contexts. The `fasttext` language identifier labels GSW (“*chuchichäschtli*”-style orthography) as standard German, so incidental dialect content cannot be reliably separated.

3.2 Annotation schema

The label space follows the `openai/privacy-filter` taxonomy: eight PII categories listed in Table 2, each encoded in a BIOES tag scheme [18] (Begin, Inside, Outside, End, Single), giving 33 classes overall (one outside class plus four BIOES positions for each of the eight categories). The byte-identical label mapping is deliberate: it lets a fine-tune started from `openai/privacy-filter`’s checkpoint reuse the pretrained classifier-head weights without remapping. The `private_` prefix is historical; all eight cells are applied

with a recall-oriented labelling policy and may flag publicly-listed institutional entities (e.g. a court switchboard phone number).

Table 2: The eight PII categories.

Category	Examples
account_number	CH-IBAN, AHV/AVS, VAT-CHE, credit cards
private_address	Bahnhofstrasse 1, 8001 Zürich
private_date	1990-01-02, 14. März 2026
private_email	alice@example.ch
private_person	Joel Barmettler, Dr. Schmidt
private_phone	+41 44 268 12 34
private_url	https://kanzlei.ch/...
secret	API keys, bearer tokens

3.3 Curation pipeline

The dataset is the output of a seven-phase pipeline. Real-text chunks are passed through (i) language detection, (ii) primary annotation by Gemma 4 26B-A4B via structured-output JSON-schema decoding, (iii) augmentation by a regex catalogue with checksum validators (ISO 13616 mod-97 for IBAN, EAN-13 mod-10 for AHV, ISO 7064 mod-11 for VAT-CHE, Luhn for credit cards), with the regex annotator taking priority on overlap. The annotated chunks are then balanced by per-document, per-value and per-cell caps to limit document-level concentration and value memorisation, then split at the document level into train, validation and test stratified by (language $\times$ source).

A separate synthetic gap-fill layer is appended to populate cells where real-Swiss-text coverage was insufficient. A slot-fill verifier pre-generates payload values (Faker_CH names, addresses and phones; checksum-validated CH-IBAN, AHV and VAT-CHE; common secret formats), then prompts Gemma to embed those exact values verbatim into natural multilingual prose (developer chat, customer-record notes, incident reports). The verifier accepts a chunk only if every payload value appears as an exact substring in the model output. Rejection rate was approximately 2 % of generations. 24,900 of these synthetic chunks are placed in the training split (22,500), validation (1,200) and test (1,200), with deterministic seeded placement and disjoint chunk IDs across splits. The remaining 146,436 chunks are real-text. Every record carries a synthetic boolean field so consumers can filter to real-only or synthetic-only as required.

3.4 Splits and held-out integrity

Documents are partitioned at the source-document level before any chunk is taken, with stratification by (language $\times$ source), so no chunk in test or validation comes from a document with a chunk in train. Synthetic chunks are placed deterministically (random seed 17). Final split sizes are 139,641 / 15,834 / 15,861 chunks for train / validation / test.

3.5 Label noise

Annotations are machine-generated; no third-party human annotators were contracted. As an upper bound on label noise we ran four independent subagent re-labellings on a stratified 176-chunk

sample and compared the dataset to the four-way majority vote, yielding an F_1 of 0.66. Inspection of disagreements indicates the gap is dominated by policy-interpretation differences (e.g. whether to label the publicly-listed switchboard number of a court as PII) rather than random error.

4 MODEL: GHEIM-CH-560M

4.1 Architecture and base-model selection

The model is a token-classification head atop xlm-roberta-large [5] (24 layers, 1024 hidden, 16 heads, 560M parameters). XLM-R-large was selected following a controlled bake-off against ZurichNLP/swissbert [23] (270M dense), summarised in Table 3. Both bases received an identical 5 $\times$ 3 hyperparameter sweep over (learning rate, layer-wise LR decay) at one epoch, after which the winning configuration per base was trained for three full epochs and selected by best validation F_1 . A third candidate, openai/privacy-filter, was deferred: the Olmo-MoE ONNX export is not currently stable in optimum [11], which would constrain JavaScript-side portability of the released model.

Table 3: Base-model bake-off.

Base model	Val F_1	Test F_1
ZurichNLP/swissbert (270M)	0.910	0.910
FacebookAI/xlm-roberta-large (560M)	0.918	0.916

In the (LR, LLRD) sweep, learning rate was the dominant factor: F_1 increased monotonically with LR over the tested range (5e-6 to 5e-5) for both bases. Layer-wise LR decay did not improve any cell at any tested value (1.0, 0.95, 0.9). The selected configuration is therefore $LR = 5 \times 10^{-5}$ without LLRD.

4.2 Training procedure

The training mix is the train split of gheim-ch-pii-171k (139,641 chunks) plus a small AI4Privacy English/email augmentation slice (approximately 14,000 chunks drawn from ai4privacy/openpii-1m [3]) added to shore up English-anchor coverage and CH-region email recall, for a total of 153,641 training chunks. The validation and test splits used for the headline metrics contain no AI4Privacy data. Hyperparameters are listed in Table 4.

Table 4: Training hyperparameters.

Hyperparameter	Value
Optimiser	AdamW [10]
Learning rate	5e-5
LR scheduler	Cosine, 5% warmup
Layer-wise LR decay	1.0 (off)
Effective batch size	128 (per-device 64 $\times$ 2 GPUs DDP)
Epochs	3; best at epoch 2.50
Max sequence length	512
Precision	bf16
Hardware	2 $\times$ NVIDIA RTX 4090
Wall time	$\approx$ 52 min train + 4 min eval

The best checkpoint was selected by `eval_overall_f1` on the validation split. The held-out test split was scored exactly once on the selected checkpoint to produce the headline numbers in Section 5.

4.3 Deployment formats

The model ships in two formats. Server-side users load `model.safetensors` (fp32 PyTorch, 2.2 GB) via `transformers` [26]; in-browser users load `onnx/model_quantized.onnx` (int8 dynamic-quantised ONNX, 552 MB) via `transformers.js` [8] on WebGPU or WebAssembly. The int8 export reproduces 99.4% of the fp32 strict F_1 (0.909 vs 0.916, -0.7 pp); the loss is concentrated in the two categories that were already weakest under fp32 (Table 5, “Quant. Δ ” column).

5 EVALUATION

5.1 Methodology

Two metrics are reported throughout. **Strict-span F_1** is the standard NER literature metric: a predicted span counts as a true positive only if its (start, end, label) tuple matches a gold span exactly. Computed via `sequeval` [14] at the token level for HF and ONNX backends, and via char-level set match for span-emitting backends (spaCy, Presidio) that lack a token grid. **Char-level F_1** (label-aware) is computed over (character index, label) pairs: it asks “did the model mask the right characters with the right category?” and is invariant to entity-segmentation policy differences across systems.

The latter matters because cross-model PII evaluation has a structural problem: different annotation conventions split the same surface form differently. `ai4privacy/openpii-1m` splits a person mention into a `GIVENNAME` and a `SURNAME` span; `gheim-ch-560m` and `CoNLL-2003` emit one combined `PERSON` span; address detectors fragment “Werdstrasse 36, 8004 Zürich” at the comma or do not. Strict-span F_1 penalises every boundary mismatch even when the underlying redaction-utility goal is met. Reporting both metrics lets the reader judge which is appropriate for their use case: strict for literature comparability, char for deployment.

5.2 Headline performance

On the held-out test split of `gheim-ch-pii-171k` (15,861 chunks), `gheim-ch-560m` attains a strict-span F_1 of 0.916 (precision 0.907, recall 0.926) and a char F_1 of 0.958. The validation/test gap is 0.002, consistent with no observable overfitting to the validation set used for checkpoint selection. The per-language $\times$ per-category breakdown is in Table 5: body cells give char F_1 per (lang, cat); the right-most columns give the per-category strict F_1 , the int8-quantisation Δ on that cell, and the gold-span support; the bottom rows give the per-language char F_1 , strict F_1 , and chunk count.

Structured categories with strong surface signals (`secret`, `email`, `phone`, `URL`) hold above 0.95 in every language; `account_number` carries almost all of the per-language variance. The Romansh `account_number` cell is the only one below 0.5: only 91 of its 117 gold spans across the test set are RM IBANs (the rest are synthetic-injected chunks), and the model under-recalls them in mixed-language contexts. Address weakness is more uniform across the four Swiss languages (0.83–0.92) and

reflects boundary placement on multi-token addresses rather than category confusion. For production use both categories are intended to be paired with a deterministic regex front-end that applies checksum validation (see `packages/gheim-py/src/gheim/detectors/composite.py`).

5.3 Comparison to other systems

Seven other open PII / NER systems were scored on the same test split under the same harness. Results are in Table 6. PER-only models (Davlan, `dslim`, `spaCy`) emit only person mentions and no other PII categories; their overall F_1 on our 8-category schema is therefore artificially low and is reported as n/a in the table; the meaningful number for those is the `private_person` cell (Table 6, last two columns).

The released model leads every other system on every cell that is directly comparable. The two metrics tell the same story: strict F_1 shows a 47-pp lead over the strongest contestant (`openai/privacy-filter`), char F_1 shows a 35-pp lead. The gap between the two metrics is largest for systems whose schema differs most from ours: `IsoTonic/distilbert_finetuned_ai4privacy_v2` gains +35 pp under char F_1 because its fine-grained schema (`FIRSTNAME`, `LASTNAME`, `MIDDLENAME`, `PREFIX`, ...) collapses under our 8-category remap and produces highly fragmented spans; `openai/privacy-filter` gains +17 pp because the Olmo tokeniser emits subword boundaries that do not align with the gold span ends.

Per-language overall char F_1 for every contestant is given in Table 7: `gheim` leads every other system in every language and is the only entry with native Romansh support. Cells where a model has no language coverage are reported as n/a (`spaCy` has no Romansh or English pipelines in this configuration; `dslim/bert-base-NER` is English-only).

5.4 Cross-domain evaluation

The released model was scored on four established external benchmarks (Table 8). For pure-NER datasets (`swissner`, `CoNLL-2003`, `WikiNeural`) only PER maps to a `gheim` category; the meaningful number is therefore the PER cell. The non-PER overall F_1 is dragged down by the model predicting categories (dates, URLs, addresses) that the NER datasets do not label.

The model’s PER F_1 stays in [0.89, 0.96] across all four datasets under both metrics. This is consistent with no overfitting to the training distribution: the four datasets cover templated synthetic English+European PII (`openpii-1m`), real Swiss news in four languages (`swissner`), the 22-year-old reference English NER benchmark (`CoNLL-2003`), and Wikipedia-grounded multilingual NER (`WikiNeural`). The notable exception is the strict-span PER score on `openpii-1m` (0.776 vs. char 0.920): `openpii-1m`’s gold splits each person mention into separate `GIVENNAME` and `SURNAME` spans, which after our remap fragment into two adjacent spans of `private_person` where our model emits one combined span. Char F_1 reports the redaction outcome (every PII character correctly masked) and is the appropriate metric for the cross-schema comparison.

On Swiss news (`swissner`, the only Swiss-news NER set that includes Romansh), per-language person-character F_1 is 0.864 (de), 0.870 (fr), 0.835 (it), 0.795 (rm).

Table 5: Held-out test split of gheim-ch-pii-171k, scored under gheim-ch-560m. Body cells are char-level F_1 for each (language, category) pair. Right-most columns give per-category marginals: strict-span F_1 (the headline NER-literature metric), the change under int8 quantisation relative to that fp32 strict cell, and the total gold-span support across all languages. Bottom rows give the per-language marginals. Smallest cell: `it_ch` $\times$ `account_number` (24 gold spans); weakest cell: `rm` $\times$ `account_number` (91 gold spans, all synthetic-injected IBANs).

Category	Char F_1 per language					Per-category marginals		
	de_ch	fr_ch	it_ch	rm	en	Strict	Quant. Δ	n_{gold}
<code>account_number</code>	0.932	0.717	0.660	0.169	0.994	0.669	-0.053	502
<code>private_address</code>	0.890	0.870	0.915	0.825	0.973	0.782	-0.021	3,507
<code>private_date</code>	0.943	0.934	0.952	0.888	0.909	0.922	-0.007	9,678
<code>private_email</code>	0.988	0.994	0.999	0.992	0.999	0.987	-0.003	1,712
<code>private_person</code>	0.938	0.948	0.962	0.897	0.951	0.915	-0.006	14,726
<code>private_phone</code>	0.989	0.985	0.993	0.995	0.997	0.986	-0.002	3,329
<code>private_url</code>	0.992	0.994	0.994	0.958	0.980	0.947	-0.002	3,991
<code>secret</code>	0.999	1.000	0.999	1.000	1.000	0.996	+0.002	1,104
Per-language char F_1	0.954	0.953	0.972	0.923	0.981	0.958		
Per-language strict F_1	0.911	0.906	0.942	0.853	0.954	0.916		
n_{chunks}	4,753	4,705	4,169	1,341	893	15,861		

Table 6: Direction A: other models on the gheim-ch-pii-171k test split. Strict and char F_1 for both the overall 8-category metric and the `private_person` cell.

Model	Strict F_1	Char F_1	PER strict	PER char
joelbarmettler/gheim-ch-560m	0.916	0.958	0.915	0.944
openai/privacy-filter (1.4B MoE, zero-shot)	0.443	0.610	0.406	0.561
Microsoft Presidio (multi-lang config)	0.434	0.562	0.442	0.521
Davlan/xlm-roberta-base-ner-hrl	n/a	n/a	0.645	0.728
Davlan/distilbert-base-multilingual-cased-ner-hrl	n/a	n/a	0.619	0.771
spaCy de/fr/it_core_news_lg	n/a	n/a	0.505	0.621
dslim/bert-base-NER (English slice)	n/a	n/a	0.512	0.770
Isotonic/distilbert_finetuned_ai4privacy_v2	0.157	0.508	0.264	0.507

Table 7: Per-language char F_1 on the gheim-ch-pii-171k test split, all eight scored systems. Rows are sorted by overall char F_1 .

Model	de_ch	fr_ch	it_ch	rm	en	Overall
joelbarmettler/gheim-ch-560m	0.954	0.953	0.972	0.923	0.981	0.958
openai/privacy-filter (1.4B MoE, zero-shot)	0.569	0.577	0.630	0.643	0.788	0.610
Microsoft Presidio (multi-lang config)	0.572	0.596	0.690	0.251	0.561	0.562
Isotonic/distilbert_finetuned_ai4privacy_v2	0.467	0.515	0.530	0.415	0.639	0.508
Davlan/distilbert-base-multilingual-cased-ner-hrl	0.350	0.340	0.368	0.362	0.282	0.348
Davlan/xlm-roberta-base-ner-hrl	0.358	0.330	0.362	0.373	0.277	0.346
spaCy de/fr/it_core_news_lg	0.308	0.298	0.337	n/a	n/a	0.314
dslim/bert-base-NER (English slice)	n/a	n/a	n/a	n/a	0.279	0.279

Table 8: Direction B: gheim-ch-560m on external benchmarks.

External benchmark	n	Overall strict	Overall char	PER strict	PER char
ai4privacy/pii-masking-openpii-1m	8,000	0.919	0.972	0.776	0.920
Babelscape WikiNeural	8,000	0.833	0.880	0.919	0.960
ZurichNLP swissner	800	0.781	0.838	0.903	0.946
CoNLL-2003	3,453	0.678	0.703	0.887	0.936

5.5 Methodology validation

To check that the harness above is sound rather than an unfair implementation, we replicated each external system on the dataset

on which it was trained and compared our reproduction to the system’s own published F_1 (Table 9). Three of four reproductions match within ± 0.03 ; the fourth (Isotonic) is discussed below.

Table 9: Methodology validation: each baseline scored on its own training-distribution dataset, against published numbers.

Replication target	Claimed	Repro	Δ
dslim/bert-base-NER on CoNLL-2003	0.910	0.913	+0.003
dslim PER on CoNLL-2003	~ 0.96	0.957	-0.003
Davlan-XLMR on WikiNeural avg	~ 0.84	0.813	-0.027
spaCy de_core on swissner_de	~ 0.83	0.841	+0.011
Isotonic distilbert on AI4Privacy	0.955	0.969 [†]	+0.014

[†]The fourth row required investigation. Isotonic’s claimed F_1 of 0.955 initially reproduced as 0.78 strict-span F_1 on the same data, an apparent 17-pp gap. The difference traces to the metric: Isotonic’s reported number is *per-token classification accuracy* (which includes the dominant O class and is much higher than per-span F_1). Reproducing Isotonic’s metric exactly (per-token accuracy including O) yields 0.969, matching their 0.955 within sampling noise. Our strict-span F_1 of 0.78 on the same data is the correct NER-literature comparison.

6 LIMITATIONS

- **Recall-oriented labelling policy.** Public Swiss documents contain extensive personally-identifiable information that is already public (court rulings, parliament records, public officials). The dataset labels these as PII and the model learns to flag them. Applications that need to distinguish public from private entities should layer a downstream public-vs-private classifier.
- **private_address test F_1 is 0.78.** Boundary placement on multi-token addresses is the dominant error mode.
- **account_number test F_1 is 0.67.** Production deployments should pair the model with a regex front-end with checksum validation (IBAN, AHV, VAT-CHE, Luhn).
- **Romansh test F_1 is 0.85,** the weakest of the five languages. The Romansh training material is dominated by a single literary/journalistic register (Rumantsch Grischun), so dialectal or technical RM text is unmeasured.
- **Swiss German dialect (GSW) is not measured.** The fast-text detector used in data preparation labels GSW as standard German.
- **Annotations are machine-generated.** The dataset was not human-annotated; the F_1 of 0.66 against a four-way subagent re-labelling is reported as an upper bound on label noise rather than as a quality claim.
- **Re-identification is not in scope.** The model is intended for redaction. It is not designed to identify or link entities to real persons and does not return entity-linked identifiers.

7 CONCLUSION

We released gheim-ch-pii-171k and gheim-ch-560m, a Swiss-tuned PII NER dataset and model targeting a deployment context (LLM-API anonymisation round-trip) that the existing open detectors handle poorly on Swiss text. The model leads every public contestant by 35-pp on the same Swiss test set under either reported metric and stays within $[0.89, 0.96]$ person-NER F_1 across four cross-domain external benchmarks. The companion dataset is released under CC BY 4.0; the model under Apache 2.0; both are on the Hugging Face Hub. The full evaluation harness, the per-system JSONs, and the reproducible build of the comparison tables are at <https://github.com/joelbarmettlerUZH/gheim>.

Future work. A privately-tuned variant covering spoken GSW, dialectal Romansh and additional account-number formats; a distilled student model in the 100M-parameter range for edge-side deployment; and a public-versus-private-entity post-classifier to relax the recall-oriented labelling policy when applications require stricter precision.

REFERENCES

- [1] David Ifeoluwa Adelani. 2022. XLM-RoBERTa-base fine-tuned on multilingual NER (high-resource languages). Hugging Face model. <https://huggingface.co/Davlan/xlm-roberta-base-ner-hrl>
- [2] AI4Privacy. 2024. PII Masking 200k: An English-language PII dataset. Hugging Face dataset. <https://huggingface.co/datasets/ai4privacy/pii-masking-200k>
- [3] AI4Privacy. 2024. PII Masking OpenPII-1M: A multilingual PII dataset. Hugging Face dataset. <https://huggingface.co/datasets/ai4privacy/pii-masking-openpii-1m>
- [4] Association Entschidsuche. 2025. Entschidsuche.ch: Open Legal Data Platform for Swiss court decisions. <https://entschidsuche.ch/docs>. Accessed: 2025-08-31.
- [5] Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. Unsupervised Cross-lingual Representation Learning at Scale. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics* (2020), 8440–8451. <https://doi.org/10.18653/v1/2020.acl-main.747>
- [6] Alejandro Hernández-Cano, Alexander Hägele, Allen Hao Huang, Angelika Romanou, Antoni-Joan Solergibert, et al. 2025. Apertus: Democratizing Open and Compliant LLMs for Global Language Environments. arXiv:2509.14233 <https://arxiv.org/abs/2509.14233>
- [7] Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. 2020. spaCy: Industrial-strength Natural Language Processing in Python. <https://doi.org/10.5281/zenodo.1212303>
- [8] Hugging Face. 2024. Transformers.js: Run Transformers directly in the browser. <https://huggingface.co/docs/transformers.js>
- [9] Benoit Jacob, Skirmantas Kligys, Bo Chen, Menglong Zhu, Matthew Tang, Andrew Howard, Hartwig Adam, and Dmitry Kalenichenko. 2018. Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference. arXiv:1712.05877 <https://arxiv.org/abs/1712.05877>
- [10] Ilya Loshchilov and Frank Hutter. 2019. Decoupled Weight Decay Regularization. arXiv:1711.05101 <https://arxiv.org/abs/1711.05101>
- [11] Microsoft. 2018. ONNX Runtime: A cross-platform inference and training accelerator. <https://onnxruntime.ai>
- [12] Microsoft. 2023. Presidio: Context-aware, pluggable and customizable PII anonymization service. <https://github.com/microsoft/presidio>
- [13] Niklas Muennighoff, Luca Soldaini, Dirk Groeneveld, Kyle Lo, Jacob Morrison, Sewon Min, Weijia Shi, Pete Walsh, Oyvind Tafjord, Nathan Lambert, Yizhong Wang, Akshita Bhagia, Mengzhou Liu, Ben Petrov, Alana Tsai, Hannaneh Hajishirzi, Noah A. Smith, Luke Zettlemoyer, Iz Beltagy, Dirk Groeneveld, Luca Soldaini, and Hannaneh Hajishirzi. 2024. OLMoE: Open Mixture-of-Experts Language Models. arXiv:2409.02060 <https://arxiv.org/abs/2409.02060>
- [14] Hiroki Nakayama. 2018. sequeval: A Python framework for sequence labeling evaluation. <https://github.com/chakki-works/sequeval>
- [15] OpenAI. 2025. Privacy Filter: A token-classification model for PII detection. Hugging Face model. <https://huggingface.co/openai/privacy-filter>
- [16] Guilherme Penedo, Hynek Kydlíček, Loubna Ben Allal, Anton Lozhkov, Margaret Mitchell, Colin Raffel, Leandro Von Werra, and Thomas Wolf. 2024.

- The FineWeb Datasets: Decanting the Web for the Finest Text Data at Scale. arXiv:2406.17557 <https://arxiv.org/abs/2406.17557>
- [17] Guilherme Penedo, Hynek Kydlíček, Vinko Sabolčec, Bettina Messmer, Negar Foroutan, Amir Hossein Kargaran, Colin Raffel, Martin Jaggi, Leandro Von Werra, and Thomas Wolf. 2025. FineWeb2: One Pipeline to Scale Them All – Adapting Pre-Training Data Processing to Every Language. *arXiv preprint arXiv:2506.20920* (2025). <https://arxiv.org/abs/2506.20920>
- [18] Lev Ratinov and Dan Roth. 2009. Design Challenges and Misconceptions in Named Entity Recognition. In *Proceedings of the Thirteenth Conference on Computational Natural Language Learning (CoNLL-2009)*. 147–155.
- [19] David Rosenthal and Livio Veraldi. 2025. Training AI language models with third-party content and data from a legal perspective. *Jusletter-IT* (March 2025). <https://doi.org/10.38023/bec9257a-a6bb-41cf-b2f3-fda8ae3b448d>
- [20] Simone Tedeschi, Valentino Maiorca, Niccolò Campolungo, Francesco Cecconi, and Roberto Navigli. 2021. WikiNEuRal: Combined Neural and Knowledge-based Silver Data Creation for Multilingual NER. In *Findings of the Association for Computational Linguistics: EMNLP 2021*. 2521–2533. <https://doi.org/10.18653/v1/2021.findings-emnlp.215>
- [21] The Swiss Parliament. 2024. Curia Vista: business database of the Swiss Federal Assembly. <https://www.parlament.ch/en/ratsbetrieb/curia-vista>. Accessed: 2024-04-24.
- [22] Erik F. Tjong Kim Sang and Fien De Meulder. 2003. Introduction to the CoNLL-2003 Shared Task: Language-Independent Named Entity Recognition. In *Proceedings of the Seventh Conference on Natural Language Learning at HLT-NAACL 2003*. 142–147.
- [23] Jannis Vamvas, Noëmi Aepli, and Rico Sennrich. 2023. SwissBERT: The Multilingual Language Model for Switzerland. arXiv:2303.13310 <https://arxiv.org/abs/2303.13310>
- [24] Jannis Vamvas, Noëmi Aepli, and Rico Sennrich. 2023. SwissNER: A NER test set for Swiss Standard German, French, Italian, and Romansh. Hugging Face dataset. <https://huggingface.co/datasets/ZurichNLP/swissner>
- [25] Jannis Vamvas, Ignacio Pérez Prat, Not Battista Soliva, others, and Rico Sennrich. 2025. Expanding the WMT24++ Benchmark with Rumantsch Grischun, Sursilvan, Sutsilvan, Surmiran, Puter, and Vallader. arXiv:2509.03148 <https://arxiv.org/abs/2509.03148>
- [26] Thomas Wolf, Lysandre Debut, Victor Sanh, et al. 2020. Transformers: State-of-the-Art Natural Language Processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. 38–45. <https://doi.org/10.18653/v1/2020.emnlp-demos.6>